

# Assignment 5

---

A. Create a directory /hadoop/hdfs/ in HDFS

- `hdfs dfs -mkdir /hadoop/hdfs`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -mkdir /hdfs
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /
Found 10 items
drwxr-xr-x  - hadoop supergroup      0 2024-09-03 18:04 /airlines
drwxr-xr-x  - hadoop supergroup      0 2024-08-31 16:54 /flights
drwxr-xr-x  - hadoop supergroup      0 2024-09-03 18:01 /flights1
drwxr-xr-x  - hadoop supergroup      0 2024-09-05 23:40 /hadoop
drwxr-xr-x  - hadoop supergroup      0 2022-11-21 15:25 /hbase
drwxr-xr-x  - hadoop supergroup      0 2024-09-05 23:45 /hdfs
drwxr-xr-x  - hadoop supergroup      0 2024-09-04 17:54 /testStreaming
drwxrwxrwx  - hadoop supergroup      0 2022-11-21 15:12 /tmp
drwxr-xr-x  - hadoop supergroup      0 2024-09-03 17:07 /user
drwxr-xr-x  - hadoop supergroup      0 2024-08-31 12:42 /wordcount
hadoop@hadoop-VirtualBox:~$
```

B. Create a temp directory in Hadoop. Run HDFS command to delete "temp" directory.

- `hdfs dfs -mkdir /temp`
- `hdfs dfs -rm /temp`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -mkdir /temp
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /
Found 11 items
drwxr-xr-x - hadoop supergroup 0 2024-09-03 18:04 /airlines
drwxr-xr-x - hadoop supergroup 0 2024-08-31 16:54 /flights
drwxr-xr-x - hadoop supergroup 0 2024-09-03 18:01 /flights1
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:40 /hadoop
drwxr-xr-x - hadoop supergroup 0 2022-11-21 15:25 /hbase
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:45 /hdfs
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:46 /temp
drwxr-xr-x - hadoop supergroup 0 2024-09-04 17:54 /testStreaming
drwxrwxrwx - hadoop supergroup 0 2022-11-21 15:12 /tmp
drwxr-xr-x - hadoop supergroup 0 2024-09-03 17:07 /user
drwxr-xr-x - hadoop supergroup 0 2024-08-31 12:42 /wordcount
hadoop@hadoop-VirtualBox:~$ hdfs dfs -rm -r /temp
Deleted /temp
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /
Found 10 items
drwxr-xr-x - hadoop supergroup 0 2024-09-03 18:04 /airlines
drwxr-xr-x - hadoop supergroup 0 2024-08-31 16:54 /flights
drwxr-xr-x - hadoop supergroup 0 2024-09-03 18:01 /flights1
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:40 /hadoop
drwxr-xr-x - hadoop supergroup 0 2022-11-21 15:25 /hbase
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:45 /hdfs
drwxr-xr-x - hadoop supergroup 0 2024-09-04 17:54 /testStreaming
drwxrwxrwx - hadoop supergroup 0 2022-11-21 15:12 /tmp
drwxr-xr-x - hadoop supergroup 0 2024-09-03 17:07 /user
drwxr-xr-x - hadoop supergroup 0 2024-08-31 12:42 /wordcount
hadoop@hadoop-VirtualBox:~$
```

C. List all the files/directories for the given HDFS destination path.

- `hdfs dfs -ls /`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /
Found 10 items
drwxr-xr-x - hadoop supergroup 0 2024-09-03 18:04 /airlines
drwxr-xr-x - hadoop supergroup 0 2024-08-31 16:54 /flights
drwxr-xr-x - hadoop supergroup 0 2024-09-03 18:01 /flights1
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:40 /hadoop
drwxr-xr-x - hadoop supergroup 0 2022-11-21 15:25 /hbase
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:45 /hdfs
drwxr-xr-x - hadoop supergroup 0 2024-09-04 17:54 /testStreaming
drwxrwxrwx - hadoop supergroup 0 2022-11-21 15:12 /tmp
drwxr-xr-x - hadoop supergroup 0 2024-09-03 17:07 /user
drwxr-xr-x - hadoop supergroup 0 2024-08-31 12:42 /wordcount
hadoop@hadoop-VirtualBox:~$
```

D. Command that will list the directories in /hadoop folder.

- `hdfs dfs -ls /hadoop`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /hadoop
Found 1 items
drwxr-xr-x - hadoop supergroup 0 2024-09-05 23:40 /hadoop/issac
hadoop@hadoop-VirtualBox:~$
```

E. Command to list recursively all files in hadoop directory and all subdirectories in hadoop directory.

- `hdfs dfs -ls -R /hadoop`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls -R /hadoop
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir1
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir2
drwxr-xr-x  - hadoop supergroup      0 2024-09-05 23:40 /hadoop/issac
hadoop@hadoop-VirtualBox:~$
```

F. List all the directories inside /hadoop/hdfs/ directory which start with 'dir'.

- `hdfs dfs -ls /hadoop/hdfsdir*`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls -d /hadoop/hdfs/dir*
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir1
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir2
hadoop@hadoop-VirtualBox:~$
```

G. Create a temp.txt file. Copy this file from the local file system to HDFS.

- `touch temp.txt`
- `hdfs dfs -put temp.txt /hadoop/hdfs`

```
hadoop@hadoop-VirtualBox:~$ gedit tmp.txt
hadoop@hadoop-VirtualBox:~$ hdfs dfs -put /home/hadoop/tmp.txt /hadoop/hdfs/
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /hadoop/hdfs
Found 3 items
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir1
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir2
-rw-r--r--  1 hadoop supergroup     25 2024-09-06 00:06 /hadoop/hdfs/tmp.txt
hadoop@hadoop-VirtualBox:~$
```

H. Copy the file from HDFS to the local file system.

- `hdfs dfs -get hadoop/hdfs/temp.txt /.`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -get /hadoop/hdfs/tmp.txt/ .
hadoop@hadoop-VirtualBox:~$ ls
Assignment4.ipynb  Desktop      Downloads    install-tl-20240830  metastore_db  Pictures  Templates  Videos
derby.log         Documents    examples.desktop  install-tl-unx.tar.gz  Music         Public    tmp.txt
hadoop@hadoop-VirtualBox:~$
```

I. Command to copy from local directory with the source being restricted to a local file reference.

- `hdfs dfs -put /temp1.txt /hadoop/hdfs`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -put ./temp1.txt /hadoop/hdfs/
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls
Found 3 items
drwxr-xr-x  - hadoop supergroup      0 2024-09-05 23:44 hdfs
drwxrwxrwx  - hadoop supergroup      0 2022-11-21 15:11 warehouse
drwxr-xr-x  - hadoop supergroup      0 2024-08-31 14:32 wordcount
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /hdfs
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /hadoop/hdfs
Found 4 items
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir1
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir2
-rw-r--r--  1 hadoop supergroup     17 2024-09-06 20:58 /hadoop/hdfs/temp1.txt
-rw-r--r--  1 hadoop supergroup     25 2024-09-06 00:06 /hadoop/hdfs/tmp.txt
```

J. Command to copy to a local directory with the source being restricted to a local file reference.

- `hdfs dfs -get /hadoop/hdfs/temp.txt ./home`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -get /hadoop/hdfs/temp.txt ./home
hadoop@hadoop-VirtualBox:~$ ls
Assignment4.ipynb  Downloads  install-tl-unx.tar.gz  Public  Videos
derby.log         examples.desktop  metastore_db          temp1.txt
Desktop          home          Music                 Templates
Documents        install-tl-20240830  Pictures              tmp.txt
hadoop@hadoop-VirtualBox:~$
```

K. Command to move from the local directory source to Hadoop directory.

- `hdfs dfs -moveFromLocal /temp2.txt /hadoop/hdfs`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -moveFromLocal ./temp2.txt /hadoop/hdfs/
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /hadoop/hdfs
Found 5 items
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir1
drwxr-xr-x  - hadoop supergroup      0 2024-09-06 00:02 /hadoop/hdfs/dir2
-rw-r--r--  1 hadoop supergroup     17 2024-09-06 20:58 /hadoop/hdfs/temp1.txt
-rw-r--r--  1 hadoop supergroup     17 2024-09-06 21:06 /hadoop/hdfs/temp2.txt
-rw-r--r--  1 hadoop supergroup     25 2024-09-06 00:06 /hadoop/hdfs/tmp.txt
hadoop@hadoop-VirtualBox:~$
```

L. Delete the directory and any content under it recursively.

- `hdfs dfs -rm -r /hadoop/hdfs`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -rm -r /hadoop/hdfs
Deleted /hadoop/hdfs
hadoop@hadoop-VirtualBox:~$
```

M. List the files and show format file sizes in a human-readable fashion.

- `hdfs dfs -du -h /wordcount`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls -h /wordcount
Found 2 items
-rw-r--r-- 1 hadoop supergroup 429.4 K 2024-08-31 12:36 /wordcount/HarryPotter_Sorcerers_Stone.txt
-rw-r--r-- 1 hadoop supergroup 1.1 M 2024-08-31 12:42 /wordcount/HarryPotterandtheDeathlyHallows.txt
hadoop@hadoop-VirtualBox:~$
```

N. Take a source file and output the file in text format on the terminal.

- `hdfs dfs -cat tmp.txt`

```
hadoop@hadoop-VirtualBox:~$ cat tmp.txt
Hi, This is a dummy file
hadoop@hadoop-VirtualBox:~$
```

O. Display the content of the HDFS file test on your /user/hadoop2 directory.

- `hdfs dfs -cat /wordcount/testHP.txt`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -put ./testHP.txt /wordcount
hadoop@hadoop-VirtualBox:~$ hdfs dfs -cat /wordcount/testHP.txt
Mr. and Mrs. Dursley, of number four, Privet Drive, were proud to say
that they were perfectly normal, thank you very much. They were the last
people you'd expect to be involved in anything strange or mysterious,
because they just didn't hold with such nonsense.
hadoop@hadoop-VirtualBox:~$
```

P. Append the content of a local file test1 to an HDFS file test2.

- `hdfs dfs -appendToFile /tmp.txt /wordcount/testHP.txt`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -appendToFile ./tmp.txt /wordcount/testHP.txt
hadoop@hadoop-VirtualBox:~$ hdfs dfs -cat /wordcount/testHP.txt
Mr. and Mrs. Dursley, of number four, Privet Drive, were proud to say
that they were perfectly normal, thank you very much. They were the last
people you'd expect to be involved in anything strange or mysterious,
because they just didn't hold with such nonsense.
Hi, This is a dummy file
hadoop@hadoop-VirtualBox:~$
```

Q. Show the capacity, free, and used space of the filesystem.

- `hdfs dfs -df`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -df
Filesystem              Size         Used    Available   Use%
hdfs://localhost:9000 41954803712 303247659 24193789952    1%
hadoop@hadoop-VirtualBox:~$
```

R. Show the capacity, free, and used space of the filesystem. Add parameter formats the sizes of files in a human-readable fashion.

- `hdfs dfs -df -h`





- `hdfs dfsadmin -safemode leave`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfsadmin -safemode leave
Safe mode is OFF
hadoop@hadoop-VirtualBox:~$
```

W. HDFS command to format NameNode.

- `hdfs namenode -format`

X. Create a file named `hdfstest.txt` and change its number of replications to 3.

- `hdfs dfs -setrep 3 /wordcount/testHP.txt`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -setrep 3 /wordcount/testHP.txt
Replication 3 set: /wordcount/testHP.txt
hadoop@hadoop-VirtualBox:~$
```

Y. Write command to display the number of replicas for `hdfstest.txt` file.

- `hdfs dfs -stat %r /wordcount/testHP.txt`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -stat %r /wordcount/testHP.txt
3
hadoop@hadoop-VirtualBox:~$
```

Z. Write command to display the status of file "`hdfstest.txt`" like block size, file size in bytes.

- `hdfs dfs -stat /wordcount/testHP.txt`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -stat /wordcount/testHP.txt
2024-09-06 15:47:28
hadoop@hadoop-VirtualBox:~$
```

AA. Write HDFS command to change file permission from `rw-r--r--` to `rw-rw-r--` for `hdfstest.txt`.

- `hdfs dfs -chmod 761 /wordcount/testHP.txt`

```
hadoop@hadoop-VirtualBox:~$ hdfs dfs -chmod 761 /wordcount/testHP.txt
hadoop@hadoop-VirtualBox:~$ hdfs dfs -ls /wordcount
Found 3 items
-rw-r--r-- 1 hadoop supergroup 439742 2024-08-31 12:36 /wordcount/HarryPotter_Sorcerers_Stone.txt
-rw-r--r-- 1 hadoop supergroup 1164038 2024-08-31 12:42 /wordcount/HarryPotterandtheDeathlyHallows.txt
-rwxrw---x 3 hadoop supergroup 288 2024-09-06 21:17 /wordcount/testHP.txt
hadoop@hadoop-VirtualBox:~$
```